

Section 2.3: Scheduling to minimize makespan

Makespan Scheduling on Parallel Machines

Input:

m machines


n jobs with processing times $p_1, p_2, \dots, p_n \in \mathbb{Z}^+$

Output:

Assignment of jobs to machines s.t. the **makespan** is minimized


time when last job finishes

Ex:



$$\text{makespan} = \max\{C_1, C_2, C_3, C_4\} = C_2$$

How could this schedule be improved?



Local Search Alg:

Repeat

job $l \leftarrow$ job that finishes last

If there is any machine i where job l would finish earlier

Move job l to machine i

Until job l is not moved

Theorem 2.5

The local search alg. is a $(2 - \frac{1}{m})$ -approx. alg.

Proof:

Lower bounds on OPT:


$$\text{OPT} \geq p_{\max} = \max_{1 \leq j \leq n} p_j,$$

because the machine i with the largest job j has $C_i \geq p_j$.

$$\text{OPT} \geq \frac{P}{m}, \text{ where } P = \sum_{j=1}^n p_j$$

Since this is the average completion time of the machines.

Upper bound on alg.'s makespan:



$P \geq m \cdot S_l + p_l$, since all machines are busy until S_l

\Downarrow

$$S_l \leq \frac{P - p_l}{m}$$

$$p_l \leq p_{\max}$$

$$\begin{aligned} C_{\max} &= S_l + p_l \\ &\leq \frac{P - p_l}{m} + p_l \\ &= \frac{P}{m} + \left(1 - \frac{1}{m}\right) p_l \\ &\leq \text{OPT} + \left(1 - \frac{1}{m}\right) \text{OPT} \\ &= \left(2 - \frac{1}{m}\right) \text{OPT} \end{aligned}$$

□

What would be a natural greedy alg.?

List Scheduling (LS)

For $j \leftarrow 1$ to n
Schedule job j on currently least loaded machine


What is the approx. ratio of LS?

What properties of the local search alg. did we use to prove $2 - \frac{1}{m}$?

We used only the fact that all machines are busy at least until S_ℓ .

Is this also true for LS?

Yes:



LS would not have placed job ℓ on machine 3.

Theorem 2.6: LS is a $(2 - \frac{1}{m})$ -approx. alg.

Note that $\frac{C_\ell}{OPT} < 2 - \frac{1}{m}$, unless $p_\ell = p_{max}$

Thus, it seems advantageous to schedule short jobs last.

Longest Processing Time (LPT)

For each job j , in order of decreasing processing times
Schedule job j on currently least loaded machine


Theorem 2.7: LPT is a $(\frac{4}{3} - \frac{1}{3m})$ -approx. alg.

Proof:

Number the jobs s.t. $p_1 \geq p_2 \geq \dots \geq p_n$.

Then the indices indicate the order in which the jobs are scheduled.

Let job l be a job to finish last:



We can assume that $l = n$:

Let $I = \{p_1, p_2, \dots, p_n\}$ and $I' = \{p_1, p_2, \dots, p_l\}$.

Then, $LPT(I) = LPT(I')$, since jobs $l+1, \dots, n$ finish no later than job l .

Moreover, $OPT(I') \leq OPT(I)$.

Thus, if we prove $LPT(I')/OPT(I') \leq \frac{4}{3}$, we have proven $LPT(I)/OPT(I) \leq \frac{4}{3}$ (since $LPT(I)/OPT(I) \leq LPT(I')/OPT(I')$).

(Or said in a different way, we can ignore the jobs $l+1, \dots, n$.)

Thus, we can assume that no job is shorter than job l . (This will be used in Case 2 below.)

Case 1: $p_k \leq \frac{1}{3} \cdot \text{OPT}$

By the proof of Thm 2.5,

$$\begin{aligned} \text{LPT} &\leq \text{OPT} + \frac{m-1}{m} p_k \leq \text{OPT} + \frac{m-1}{m} \cdot \frac{1}{3} \cdot \text{OPT} \\ &= \left(\frac{4}{3} - \frac{1}{3m}\right) \text{OPT} \end{aligned}$$

Case 2: $p_k > \frac{1}{3} \cdot \text{OPT}$

In this case, all jobs are longer than $\frac{1}{3} \cdot \text{OPT}$.
Hence, in OPT's schedule, each machine has ≤ 2 jobs, i.e., $n \leq 2m$.

In this case, $\text{LPT} = \text{OPT}$:

p_1	
p_2	
p_3	p_8
p_4	p_7
p_5	p_6

Proof of this claim:
Exercise 2.2

□

From the proof of Thm 2.7 we learned:

If job l is longer than $\frac{1}{3} \cdot \text{OPT}$, then $\text{LPT} = \text{OPT}$.

Otherwise, $\text{LPT} \leq \text{OPT} + p_l \leq \frac{4}{3} \cdot \text{OPT}$.

(Recall that job l is the job to finish last.)

Could we balance the two cases better?

What if we first schedule all jobs of length $\geq \frac{1}{4} \cdot \text{OPT}$ optimally, and then use LPT for the remaining jobs?

What would the approximation ratio be?

Does the schedule of the long jobs have to be optimal?

Section 3.2: Makespan Scheduling - A PTAS

Idea for PTAS:

Partition the jobs into two sets (long and short jobs):

$> \varepsilon \cdot \text{OPT}$	$\leq \varepsilon \cdot \text{OPT}$
$\underbrace{p_1, p_2, \dots, p_x}_{\text{Schedule using rounding and dyn. prg. as for bin packing}}$	$\underbrace{p_{x+1}, \dots, p_n}_{\text{Then use LPT}}$

We will derive a family of algorithms with an algorithm, B_k , for each $k \in \mathbb{Z}^+$. ($\varepsilon = \frac{1}{k}$)

Let $P = \sum_{i=1}^n p_i$ (as before).

Job j is short, if $p_j \leq \frac{P}{km}$, i.e., if it is at most $\frac{1}{k}$ of the average machine load.

Otherwise, it is long.

The alg. will be poly. in m , but not in k . Thus, the algorithm will be a PTAS, not an FPTAS.

#long jobs $< km$

Hence, #schedules of long jobs $< m^{km}$

(choose one of m machines for each job).

Thus, if $k, m \in O(1)$, we can find an optimal schedule for the long jobs in time $O(1)$.

Otherwise, we can round job sizes and do dyn. prg. as for the bin packing problem:

Scheduling the long jobs:

- (1) „Guess“ an optimal makespan T
- (2) Round down each job size to the nearest multiple of T/k^2 .
- (3) Use dyn. prg. to check whether optimal makespan $\leq T$ for rounded long jobs.

Do binary search for T on the interval $[L, U]$, where

$$L = \max \left\{ \left\lceil \frac{P}{m} \right\rceil, p_{\max} \right\}$$

$$U = \left\lfloor \frac{P - p_{\max}}{m} + p_{\max} \right\rfloor = \left\lfloor \frac{P + (m-1)p_{\max}}{m} \right\rfloor$$